

BRAMBLELOOP FINAL BUILD v1.1

PRIORITY ZERO AUTONOMY / PARALLEL COMPLETION / STORE FOUNDATION

The previous Final Build run ended after approximately 17 hours. Preserve everything it accomplished.

I am attaching Brambleloop_FINAL_Master_v1.1_Owner_Command_Center.pdf. This is now the CANONICAL FINAL MASTER SPECIFICATION.

Do NOT restart Final Build from scratch. Reconcile v1.1 as a delta against actual current repository reality and continue from there.

1. OWNER RULING - OA-0003 IS RESOLVED

Rented/cloud always-on hosting is authorized and REQUIRED.

Brambleloop must not depend on my PC being powered on, Windows being logged in, my PC not sleeping, Claude, Codex, Fable, ChatGPT, or any development conversation remaining open.

The Windows watchdog/scheduled-task solution is useful development fallback evidence, but it is NOT the production solution. M-0019 must proceed.

Prefer the simplest durable always-on solution consistent with the Master and existing Railway/cloud architecture. Identify any new recurring cost before incurring it.

2. PRIORITY ZERO - FOUNDATIONAL RULE #1 IS CURRENTLY VIOLATED

This is the most important instruction in this prompt.

24/7 autonomous operation and continuous self-improvement were founding requirements, NOT new v1.1 features. The previous run proved the factory can stop when my PC sleeps or loses its logged-in user session. Treat that as a CRITICAL FINAL BUILD DEFECT.

Claude, Codex and Fable are development workers. THEY ARE NOT BRAMBLELOOP'S PRODUCTION BRAIN.

Brambleloop itself must continuously OBSERVE -> determine useful work -> prioritize -> create missions/jobs -> assign specialist agents -> execute -> independently validate -> consume outputs downstream -> create follow-on work -> measure outcomes -> identify weaknesses/opportunities -> propose improvements -> independently evaluate -> sandbox/test -> promote or reject -> monitor -> rollback when worse -> persist lessons -> alter future behaviour -> repeat.

AN EMPTY QUEUE MUST NOT MEAN AN IDLE COMPANY. A heartbeat, scheduler, or module name is not proof. I want useful autonomous company activity.

3. AUDIT EVERY EXISTING PRODUCTION AGENT / DEPARTMENT

Before inventing duplicate architecture, discover what already exists.

Persist a machine-readable autonomy map: department/agent -> mission -> authority -> wake mechanism -> cadence/event -> evidence inputs -> autonomous work-generation -> mission creation -> execution -> durable output -> downstream consumer -> follow-on work -> learning signal -> improvement path -> next wake -> blocker.

Audit Executive/COO, Intelligence, Product/Design, Product Truth/QA, Visual, Store/Commerce, Customer Support, Finance/Accounting, Growth/Marketing, Learn/Improvement, and Platform/Reliability.

Any broken link is a defect. Repair executable defects; do not merely document them.

4. CONTINUOUS SELF-IMPROVEMENT IS NON-NEGOTIABLE

Verify Learn/Improve actually changes future behaviour through measured evidence.

Required loop: measure -> diagnose -> propose challenger -> independent evaluation -> sandbox -> promote -> execute -> measure post-change outcome -> retain or rollback -> persist lesson -> affect future decisions.

Require real consumers and real effects across Visual, Search, Design, Support, Product Truth, Commerce, Growth, Finance and agent/tool/process performance.

Self-improvement may NEVER silently weaken Product Truth, customer truth/safety, accounting truth, authorization, security, spend controls or evidence requirements. Improvements require evaluation, monitoring and rollback.

5. MULTI-AGENT FINAL BUILD - DO NOT DO THIS SEQUENTIALLY

The previous prompt consumed approximately 17 hours. This phase MUST use bounded parallel implementation agents wherever file ownership and semantic dependencies permit.

Claude/Opus is the single integrator and certification authority. Build a dependency/ownership graph and launch the maximum sensible number of PARALLEL NON-OVERLAPPING agents.

Attempt independent lanes for: A cloud autonomy/factory/orchestration; B Learn/self-improvement; C Owner Command Center backend/API/security; D mobile-first web/PWA frontend; E controller-grade Accountant; F Store Foundation/Etsy setup/preview; G SEO/keywords/taxonomy/search intelligence; H Growth/ads/economics; I reliability/observability/failover/SLOs; J adversarial certification/autonomy evidence.

Adjust lanes for safe file ownership. Each worker gets requirement IDs, defect IDs, owned files, forbidden/shared files, interfaces, acceptance tests, runtime proof and commit/handoff format.

Do not wait for one independent lane to finish before another works. Reuse freed agent slots on the next READY independent cluster.

6. RECONCILE CODEX WORK - DO NOT REDO IT

Inspect relevant codex/* branches and durable handoffs before duplicating work.

Known checkpoints: codex/build2-assist-01 @ 19180b3 and codex/visual-v2-rnd @ 035ff0c62a42dd5fd0b616998dcf98833b5c1097.

Codex Visual V2 retained B+C = deterministic product authority/rendering + protected photoreal composition. Do not restart that research. Evidence decides reuse, supersession or rejection. Preserve provenance and discover newer Codex handoffs too.

7. FINAL MASTER v1.1 DELTA

Parse and reconcile the attached canonical v1.1 completely, especially F-880 through F-930.

Reuse existing implementations where they genuinely satisfy the requirements. Do not build duplicates merely because requirement numbers are new.

v1.1 explicitly covers cloud-resident operation, no development-chat dependency, Owner Command Center/PWA, approvals/security, 24/7 departments, executive orchestration, controller-grade accounting, persistent memory, failover, autonomy SLOs, Store Preview, company timeline, Ask Company and bounded autonomy.

8. OWNER COMMAND CENTER - THIS MUST BECOME HOW I RUN THE COMPANY

I do not require a native Android/iOS app. Build a secure, beautiful, responsive, mobile-first hosted web application/PWA available 24/7 from my phone.

Normal operation after launch must not require GitHub, Railway's dashboard, terminal/CLI, direct database access, Claude or Codex.

HOME: morning/overnight brief, revenue/profit, store status, launches, completed/in-progress work, autonomy health, incidents, opportunities and owner decisions.

APPROVALS: proposal, evidence, recommendation, uncertainty, upside/downside, max spend, reversibility, deadline, consequence of no action, approve/reject.

STORE: products/releases, listing state, Etsy read-back, imagery, pattern/PDF, price, taxonomy, SEO, performance and customer-facing preview.

MONEY: revenue, fees, refunds, tax/reserves, operating/model/ad spend, contribution/profit, cash commitments, payouts, reconciliation and forecasts with uncertainty.

OPERATIONS: departments/agents, jobs/queue/cadences, incidents/blockers, recovery/SLOs/dependencies.

AUTONOMY/LEARN: autonomous jobs generated/completed, hours since last owner/development prompt, active/blocked departments, inter-agent handoffs, experiments, improvements proposed/tested/promoted/rejected/rolled back, lessons, last useful action and next planned action.

INSIGHTS, NOTIFICATIONS and ACCOUNT must expose the evidence, controls, connected services, scopes, providers, budgets, authorities, policies, security and emergency controls needed to run the company.

ASK COMPANY must answer questions like What did Brambleloop do overnight?, Why did profit fall?, What is blocking launch?, and Why does this product need my approval?, grounded in durable evidence with provenance.

9. ACCOUNTANT - TOP-TIER CONTROLLER STANDARD

This is a serious subsystem, not a dashboard calculator.

Build/reconcile a dedicated Finance/Accounting specialist with source-linked accounting; double-entry or equivalently rigorous ledger; Etsy/order/payment/refund/fee reconciliation; correction entries; actual vs estimated costs separated; UNKNOWN never displayed as measured \$0.00; product/release/channel profitability; attributable model/API/image/ad/software/testing/benchmark costs; cash obligations; expected payouts; tax reserve; safe discretionary budget; anomaly detection; repeatable month-end close; tax/accountant-ready evidence packs; and human-accountant handoff.

Growth and Finance must be able to disagree. Growth cannot scale spend because ROAS looks good if Finance proves contribution economics or cash constraints are bad.

The Accountant cannot invent missing data, file taxes, move money, change banking, borrow, sign contracts or make regulated professional representations without applicable authority/human review.

10. STORE FOUNDATION - I WANT TO SEE THE REAL STORE SOON

Do not wait for all Launch-0 products before preparing the storefront. Treat the Etsy store itself as a product.

Bring every store-level surface that can safely be prepared without live listings to OWNER STORE PREVIEW quality: brand identity, icon/logo, banner, title/tagline, announcement, About/story, policies, FAQs, support/contact presentation, sections/category architecture, digital-product disclosures, brand voice, visual consistency, trust presentation, settings, Search Visibility readiness, customer-service readiness and owner-preview experience.

The standard is NOT fields populated. The standard is: DOES BRAMBLELOOP LOOK LIKE A POLISHED, PREMIUM, TRUSTWORTHY CROCHET-PATTERN COMPANY THAT HAS EXISTED FOR YEARS?

Make it beautiful. Mobile appearance matters. Do not publicly activate/publish without existing owner authority.

11. SEO / KEYWORDS / SEARCH - OPERATE IT CONTINUOUSLY

Preserve and implement query matching vs ranking, query-family research, coverage matrix, deepest truthful categories, attributes, all 13 tags where applicable, long-tail phrases, readable buyer-first titles, natural description relevance, hero/thumbnail CTR, Search Visibility findings, impressions, clicks, favorites, carts/orders, conversion, customer-service trust and dated Etsy policy/search changes.

Search is not a one-time launch task. Brambleloop must continuously learn from real search/conversion evidence and improve future decisions. Never keyword-stuff or make unsupported claims.

12. ADS / GROWTH

Advertising must be evidence-driven and economics-aware. Growth may research, propose, prepare, measure and optimize within authority.

Protected spend must revalidate owner authority, budget, listing eligibility, Product Truth, current trust, organic evidence where required, and Accountant contribution/cash constraints.

Measure actual business economics, not vanity metrics alone. Do not launch paid campaigns without applicable authority.

13. VISUAL OWNER RULING

Photorealistic Visual V2 is NOT intrinsically required for first launch. TRUTHFUL CUSTOMER-READY VISUAL EVIDENCE IS LAUNCH-CRITICAL.

Rendered/deterministic imagery may launch only when the actual customer-facing set passes Product Truth, stitch/fabric identity, dimensions/features/construction consistency, image QA, mobile/thumbnail-safe layout, disclosure, marketplace/publishing and customer-truth gates.

Disclosure alone cannot make inaccurate imagery acceptable. Do not weaken Product Truth or Visual thresholds. Preserve B+C as the preferred Visual V2 architecture.

14. MASTER SPEC INTEGRITY / ERRATA

Resolve reported anomalies: F-631 through F-650 reportedly undefined; F-514 through F-520 reportedly duplicated with different meanings; numbering reaches F-879 while a prior registry reported 866 records.

Do not invent requirements or silently discard duplicates. Determine parsing, numbering, historical-integration, missing-content or duplicate-ID causes. Persist MASTER_SPEC_ERRATA with evidence and canonical resolution where determinable.

15. PERMANENT PROOF STANDARD

Where applicable: requirement -> producer -> durable evidence/state -> runtime consumer -> decision -> protected action/effect -> observable result -> independent proof.

Code existence != execution. Handler registration != execution. Reading without consumption != completion. Fixture-only evidence != end-to-end proof. UNKNOWN != zero. Proxy != measurement. Documentation != runtime evidence. Green tests are necessary but not sufficient.

Never mark requirements complete merely to make counts green.

16. PRODUCTION AUTONOMY / LIGHTS-OUT ACCEPTANCE

The architecture must support a genuine lights-out test: owner PC OFF; no logged-in Windows session; Claude/Codex/Fable closed; no owner prompts generating normal work.

Brambleloop must remain hosted and alive, and evidence must show useful autonomous work rather than uptime: missions generated, independent departments working, inter-agent handoffs, follow-on work, monitoring/recovery, unrelated work continuing around gates, improvements proposed/tested where evidence permits, budgets respected, no duplicated external effects, durable restart state.

Do not fake a long soak by manipulating clocks or inventing evidence. If full duration cannot finish in this session, start the real soak when ready, persist evidence, and classify remaining duration honestly as an active certification gate.

17. CLOUD HOSTING / DEPLOYMENT PREPARATION

Because OA-0003 is resolved, prepare the production-autonomy hosting path. Prefer existing Railway/cloud infrastructure if it safely provides durable execution, scheduler/workers, durable DB/state, restart/recovery, secrets, health checks, logs, cost visibility and rollback.

Do not deploy merely because hosting is prepared. No production merge/deployment is authorized by this prompt unless separately approved.

18. PRESERVE THE PREVIOUS 17-HOUR RUN

Before substantial new integration, preserve exact repository reality: latest branch/head, RC/freeze state, local vs pushed commits, signing status, shadow rehearsal evidence, full-suite results, defect ledger, closure matrix, worker branches and durable handoff.

Do not rewrite or repeat valid work because v1.1 is attached. If commits remain local/unsigned, preserve the exact tested source and use the normal signing/push process available in the new session without altering source merely to obtain signatures.

19. STORE SETUP AND LAUNCH-0

Store Foundation and Store Preview should progress in parallel with product/listing completion.

Do not lower an existing launch-readiness product-count threshold merely to fit current inventory. If the canonical Master requires genuinely listable products, satisfy it through genuinely distinct validated products or surface a real owner decision. Do not count sizes as separate products unless the Master explicitly defines them that way.

20. AUTHORITY BOUNDARIES

Continue engineering, local/shadow testing, read-only verification and preparation.

This prompt does NOT authorize production deployment; merging an RC into production; public Etsy activation/publication; advertising spend; banking/KYC changes; tax filing; moving money; signing contracts; or irreversible external actions.

Existing narrow owner grants remain governed by their own recorded scope and expiry. Do not broaden them by inference.

21. EXECUTION DISCIPLINE

Do not stop after an audit, architecture report, one worker wave, or a green focused suite.

Integrate validated parallel work continuously. After each material integration: run focused tests, verify runtime reachability/consumption, update defects/closure, commit, push, and continue to the next READY independent work.

Owner/data/external gates must not idle unrelated work.

Use durable repository state as shared memory. Conversation context is disposable. Before session pressure becomes dangerous, collect worker results, preserve partial work, update assignments/defects/closure, commit/push and write deterministic resume instructions.

22. RELEASE / CERTIFICATION STANDARD

When all executable launch-critical v1.1 work is integrated, freeze a candidate and perform adversarial certification.

Certification must include a clean-state full suite, Final Master closure audit, adversarial COMPLETE sampling, runtime reachability/consumption audit, security/customer-truth review, finance truth checks, restart/recovery tests, cloud autonomy evidence, shadow business-chain rehearsal, store-preview review, and lights-out autonomy evidence/gate.

Do not call Final Build complete simply because tests are green.

23. STOPPING CONDITION

Continue autonomously until one of these is true:

A. All launch-critical executable Final Master v1.1 work is implemented/integrated/proven and a successor release candidate is frozen and under certification.

B. Every remaining launch-critical item is genuinely owner/data/external gated and no independent executable work remains.

C. Session limits force a durable handoff.

If A is reached, begin adversarial certification immediately rather than waiting for another prompt.

If C approaches, checkpoint and push BEFORE the limit kills workers.

24. REQUIRED END-OF-RUN REPORT

When this run finally stops, report compactly but completely:

- starting and ending SHA
- branches/workers launched and what each delivered
- requirements/defects closed, still open and gated
- proof that Priority Zero autonomy is fixed or exact remaining blocker
- cloud hosting status and cost/owner action if any
- self-improvement loops proven end-to-end
- Owner Command Center status and URL/preview if safely available
- Accountant status
- Store Foundation status and what I can preview now
- SEO/search and ads readiness
- full-suite/certification results
- lights-out soak status
- exact owner actions required next
- what remains before a real store can launch.

Do not hide uncertainty or substitute prose for evidence.

FINAL DIRECTIVE

Do not spend another day with one agent doing work that can safely be parallelized.

Do not treat Brambleloop's 24/7 autonomy as a future feature. It was Rule #1 and must be repaired/proven now.

Do not make Claude itself the company.

Build the autonomous company, make the store beautiful, give me a phone-first owner command center, give Finance a controller-grade Accountant, keep specialist agents working continuously, and prove the system keeps working and improving when every development chat is closed.

BEGIN NOW FROM CURRENT REPOSITORY REALITY.